



01 / Analytics platform

Operational data → governed warehouse → dashboards and natural-language analysis



02 / How the analytics agent answers

LLM-guided analysis loop prescribed by the agent skills; these are not fixed graph nodes.



OPERATOR SCOPE


The warehouse query sees activity across users. Caller-based row scoping is not implemented. The access middleware runs before model and tool execution.

DEGRADED DATA IS VISIBLE

Unavailable metadata sources produce notices. Empty results, errors and suspicious numbers trigger inspection or narrower queries; the model must not invent an answer.

03 / From verified analysis to a saved chart

Standard charts run inside create_chart. Custom Python figures join the same persistence path.



PERSISTENCE CONTRACT


Chart envelopes store figure/table/KPI payloads plus metadata and SQL when supplied. Artifacts are user-scoped; scratch CSV and Python figure files are intermediate inputs.

FAILURES ARE REPORTED

Rejected shell commands do not execute. Invalid figures, empty queries, oversized chart results or failed storage writes return errors instead of a successful chart path.

04 / Analytics on Google Cloud

PROPOSED HOSTING · existing analytics tools, Google Cloud infrastructure · no deployment changes made



Deployment scope


Analytics hosting view. Existing app runtime dependencies and full networking are omitted. Separate metadata databases for Dagster/Superset, ingress, IAM, secrets and backups need provisioning.

Compatibility to verify

The current object client is S3Store. Cloud Storage requires XML API interoperability, endpoint/HMAC configuration and integration tests. ClickHouse stays ClickHouse; BigQuery would require SQL and client changes.

05 / Analytics on AWS

PROPOSED HOSTING · existing analytics tools, AWS infrastructure · no deployment changes made



Deployment scope


Analytics hosting view. Existing app runtime dependencies and full networking are omitted. Separate metadata databases for Dagster/Superset, ingress, IAM, secrets and backups need provisioning.

Compatibility to verify

The current S3Store client targets S3. Configure the bucket, region, endpoint and credentials; verify access and encryption policies. ClickHouse runs on EC2-backed EKS nodes with EBS storage. Validate backup and recovery.

06 / Analytics on Microsoft Azure

PROPOSED HOSTING · existing analytics tools, Microsoft Azure infrastructure · no deployment changes made



Deployment scope


Analytics hosting view. Existing app runtime dependencies and full networking are omitted. Separate metadata databases for Dagster/Superset, ingress, IAM, secrets and backups need provisioning.

Compatibility to verify

Blob Storage is not a drop-in endpoint for the current S3Store client. Add AzureStore selection and Azure credentials, then test file operations and caller scoping. ClickHouse uses Azure Disk volumes; backup and recovery need configuration.

07 / Analytics on Oracle Cloud (OCI)

PROPOSED HOSTING · existing analytics tools, Oracle Cloud (OCI) infrastructure · no deployment changes made



Deployment scope

Analytics hosting view. Existing app runtime dependencies and full networking are omitted. Separate metadata databases for Dagster/Superset, ingress, IAM, secrets and backups need provisioning.

Compatibility to verify

Object Storage: configure the S3-compatible endpoint and customer secret keys; test S3Store file operations. Repo uses PostgreSQL 18; OCI documentation lists 14–17. Validate managed-service compatibility or self-manage PG18. Enable and verify pgvector.